

Nonstationary Extremes, II

Heteroscedastic Extremes in Practice

Manuele Leonelli

First LARS-IASC School on Computational Statistics and Data Science
Federal University of Bahia, Salvador, Brazil
November 15th, 2018



University
of Glasgow

Plan for this session

- ① Overview of available software
- ② Implement univariate methods in R:
 - (a) Block maxima approach
 - (b) r-largest observations approach
 - (c) Peaks over threshold
 - (d) Heteroscedastic extensions
- ③ Discuss assumptions, diagnostics and model selection criteria

All code to implement the methods is available at

github.com/manueleleonelli.

Why R?

Three main reasons:

- R is free!
- R is open-source
- Largest set of functions for extreme value analysis (EVA)

Other software is available:

- In Matlab: EVIM and WAFO
- GUI software: Xtremes and EXTREMES
- Specialized software: HYFRAN, GLSNet

E. Gilleland, M. Ribatet, A.G. Stephenson (2013) A software review for extreme value analysis. *Extremes* 16:103-119.

R packages

Most functions for EVA are stored in R packages to perform a wide array of analyses:

- Univariate: `ismev`, `extRemes`, `evmix`, `POT` ...
- Bayesian: `evdbayes`, `revdbayes`, `MCMC4extremes` ...
- Multivariate: `evd`, `copula`, `extremis`, `texmex` ...
- Spatial: `spatialExtremes`, `spatialADAI` ...

See [this guide](#).


All these require at least some programming, as usual in R. One notable exception is `in2extremes`.

Let's get started!

```
> library(ismev)
> data(fremantle)
> head(fremantle)
```


	Year	SeaLevel	SOI
1	1897	1.58	-0.67
2	1898	1.71	0.57
3	1899	1.40	0.16
4	1900	1.34	-0.65
5	1901	1.43	0.06
7	1903	1.19	0.47

```
> plot(fremantle[,1:2], pch=19, cex=0.5)
```



Southern Oscillation Index (SOI)

```
> ggplot(fremantle, aes(x=SOI, y=SeaLevel,  
+ color=Year)) + geom_point()
```



Fitting GEV distributions

```
> library(extRemes)
> ?fevd
```

Let's look at the **help**

Another possibility is to use the **ismev** package.

E. Gilleland, R.W. Katz (2016) extRemes 2.0: An extreme value analysis package in R. *Journal of Statistical Software* 72:1-39.

```
> fremantle_fit <- fevd(fremantle$SeaLevel,type="GEV")
> summary(fremantle_fit)

fevd(x = fremantle$SeaLevel, type = "GEV")

[1] "Estimation Method used: MLE"
```

Negative Log-Likelihood Value: -43.56663

Estimated parameters:


location	scale	shape
1.4823417	0.1412723	-0.2174282

Standard Error Estimates:

location	scale	shape
0.01672527	0.01149706	0.06378114


Diagnostics: qq-plot

```
> plot(fremantle_fit,type=c("qq"),main="")
```



Diagnostics: densities


```
> plot(fremantle_fit,type=c("density"),main="")
```



N = 86 Bandwidth = 0.04409

Diagnostics: return levels

```
> plot(fremantle_fit,type=c("rl"),main="")
```



Inference: confidence intervals

```
> ci(fremantle_fit,type="parameter")  
fevd(x = fremantle$SeaLevel, type = "GEV")
```

```
[1] "Normal Approx."
```

	95% lower CI	Estimate	95% upper CI
location	1.4495608	1.4823417	1.51512265
scale	0.1187385	0.1412723	0.16380615
shape	-0.3424370	-0.2174282	-0.09241948

Inference: return levels

```
> return.level(fremantle_fit, return.period=
+             c(2,10,20,50,100,500), do.ci=TRUE)
fevd(x = fremantle$SeaLevel, type = "GEV")
```

```
[1] "Normal Approx."
```

	95% lower CI	Estimate	95% upper CI
2-year return level	1.498259	1.532110	1.565962
10-year return level	1.689876	1.733753	1.777631
20-year return level	1.739108	1.791463	1.843817
50-year return level	1.785590	1.853927	1.922265
100-year return level	1.810194	1.893106	1.976017
500-year return level	1.843731	1.963815	2.083900

Bayesian GEV fitting

Let's look at the **help** again.


```
> fremantle_fit_Bayes <- fevd(fremantle$SeaLevel,  
+                             type="GEV",method="Bayesian")  
> summary(fremantle_fit_Bayes)
```

```
[1] "Quantiles of MCMC Sample from Posterior Distribution"
```

	2.5% Posterior	Mean	97.5%
location	1.4432050	1.4806897	1.51878531
scale	0.1225513	0.1451410	0.17234688
shape	-0.3333519	-0.2014116	-0.05050336

Bayesian return levels


```
> plot(fremantle_fit_Bayes,type="rl",main="")
```




Convergence checking

```
> plot(fremantle_fit_Bayes,type="trace")
```


```
fevd(x = fremantle$SeaLevel, type = "GEV", method = "Bayesian")
```






N = 9500 Bandwidth = 0.002669



N = 9500 Bandwidth = 0.001813




N = 9500 Bandwidth = 0.009979



Estimated location of maximum Sea Level

```
> plot(fremantle[,1:2],pch=19,cex=0.5)  
> abline(h=fremantle_fit$results$par[1],lwd=2)
```



Non-stationary fitting

```
> fremantle$Time <- fremantle$Year-min(fremantle$Year)
> fremantle_fit_time <- fevd(SeaLevel,data=fremantle,
+                             location.fun = ~Time)
> summary(fremantle_fit_time)
```

Estimated parameters:

mu0	mu1	scale	shape
1.382225781	0.002032151	0.124326256	-0.125309664

Standard Error Estimates:

mu0	mu1	scale	shape
0.0288458801	0.0004982185	0.0103749710	0.0682389684

Is there a time effect?


```
> ci(fremantle_fit_time,type="parameter")  
fevd(x = SeaLevel, data = fremantle, location.fun = ~Time)  
  
[1] "Normal Approx."
```

	95% lower CI	Estimate	95% upper CI
mu0	1.32568889	1.382225781	1.438762667
mu1	0.00105566	0.002032151	0.003008641
scale	0.10399169	0.124326256	0.144660825
shape	-0.25905558	-0.125309664	0.008436257

Time-varying return levels


```
> plot(fremantle_fit_time, "rl", rperiods=c(5, 10, 50))
```

fevd(x = SeaLevel, data = fremantle, location.fun = ~Time)



Time-varying location of Sea Level

```
> plot(fremantle[,1:2],pch=19,cex=0.5)  
> abline(a=fremantle_fit_time$results$par[1],  
+        b=fremantle_fit_time$results$par[2])
```




Model selection via hypothesis testing

```
> dev <- 2*(-fremantle_fit_time$results$value  
+  
+fremantle_fit$results$value)  
> dev > qchisq(0.95,1)  
[1] TRUE  
> 1-pchisq(dev,1)  
[1] 0.0003671508
```

Model	AIC	BIC
Stationary	-81.1	-73.8
Non-Stationary	-91.8	-82.0

A quadratic time effect?



Deviance test has p-value 0.2231935 - *No significant quadratic effect*

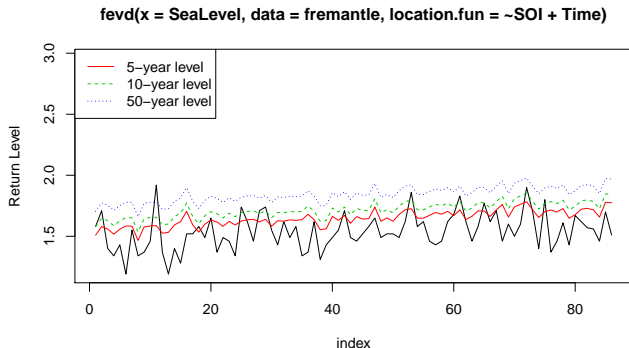
Adding SOI to the model

```
> fremantle_fit_SOI <- fevd(SeaLevel, data=fremantle,  
+                           location.fun = ~SOI+Time)  
> ci(fremantle_fit_SOI, type="parameter")  
fevd(x = SeaLevel, data = fremantle, location.fun = ~SOI + Time)  
[1] "Normal Approx."
```

	95% lower CI	Estimate	95% upper CI
mu0	1.328079958	1.384328499	1.44057704
mu1	0.016019518	0.054519871	0.09302022
mu2	0.001141007	0.002114038	0.00308707
scale	0.101075964	0.120733366	0.14039077
shape	-0.276802887	-0.149992773	-0.02318266

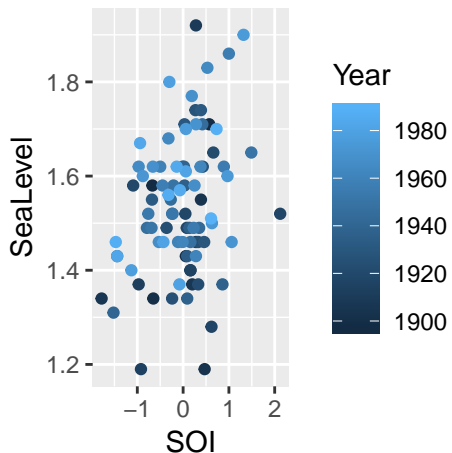
Time and SOI dependent returns

```
> plot(fremantle_fit_SOI, type="rl", rperiod=c(5, 10, 50))
```



Deviance test has p-value 5.47286e-06 - *Significant SOI effect*

What about an interaction?



We can test this easily using `location.fun = SOI+Time`.

However the deviance test has p-value 0.6162236, thus no interaction effect.

Extracting maxima


```
> data(Fort)
```

```
> head(Fort)
```

	obs	tobs	month	day	year	Prec
1	1	1	1	1	1900	0
2	2	2	1	2	1900	0
3	3	3	1	3	1900	0
4	4	4	1	4	1900	0
5	5	5	1	5	1900	0
6	6	6	1	6	1900	0

Extracting maxima

```
> bmFort <- blockmaxxer(Fort,  
+      blocks = Fort$year, which="Prec")
```




Venice sea levels

```
> data(venice)
```

```
> head(venice)
```

	Year	r1	r2	r3	r4	r5	r6	r7	r8	r9	r10
1	1931	103	99	98	96	94	89	86	85	84	79
2	1932	78	78	74	73	73	72	71	70	70	69
3	1933	121	113	106	105	102	89	89	88	86	85
4	1934	116	113	91	91	91	89	88	88	86	81
5	1935	115	107	105	101	93	91	NA	NA	NA	NA
6	1936	147	106	93	90	87	87	87	84	82	81

Venice sea levels




Fitting r-largest observations

```
> venice <- venice[-5,] #omit NA
> venice$time <- seq(1:nrow(venice)) #add time variable
> venice_fit_time <- fevd(r1,type="GEV",
+                          data = venice,location.fun = ~time)
> # Using the package ismev
> venice_fit_time_5 <- rlarg.fit(venice[,2:6],
+                                ydat=venice,mul=12)
> venice_fit_time_10 <- rlarg.fit(venice[,2:11],
+                                 ydat=venice, mul=12)
```

r	β_0	β_1	σ	ξ
1	96.7 (4.3)	0.59 (0.14)	14.6 (1.6)	-0.022 (0.084)
5	104.1 (2.0)	0.47 (0.06)	12.3 (0.8)	-0.033 (0.043)
10	104.54 (1.7)	0.49 (0.04)	11.75 (0.7)	-0.06 (0.028)

Time-dependent location


Various diagnostics can be plotted using `rlarg.diag`



Perhaps a “seasonal” effect?


A financial application: the Dow Jones Index

```
> data(dowjones)
> plot(dowjones, type='l')
```



The negative log-returns

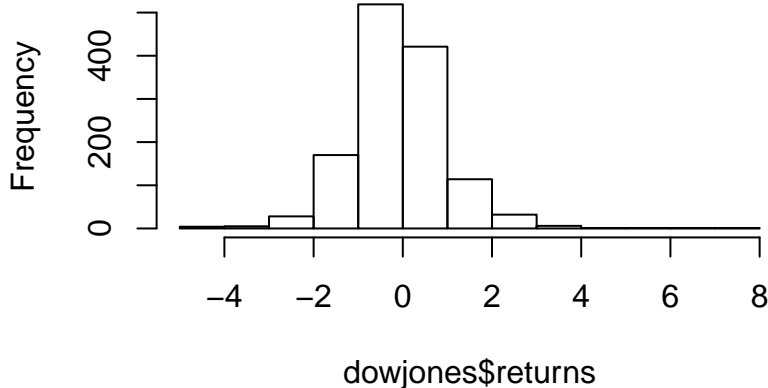
```
> dowjones <- data.frame(dowjones[-1,],  
+                         returns=-100*diff(log(dowjones$Index)))  
> plot(dowjones$Date,dowjones$returns,type="l")
```



The returns' distribution


```
> hist(dowjones$returns)
```

Histogram of dowjones\$returns



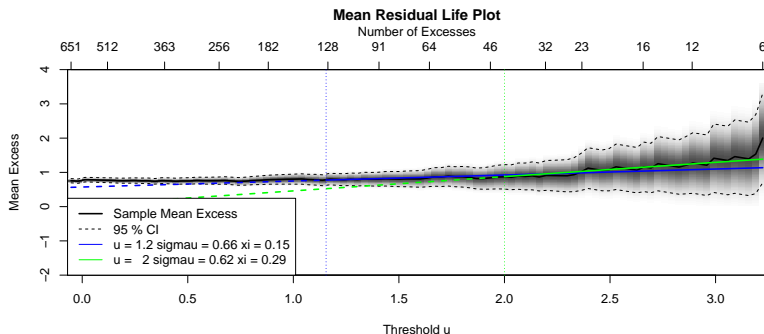
Choosing the threshold: MRL plot

```
> mrl.plot(dowjones$returns)
```



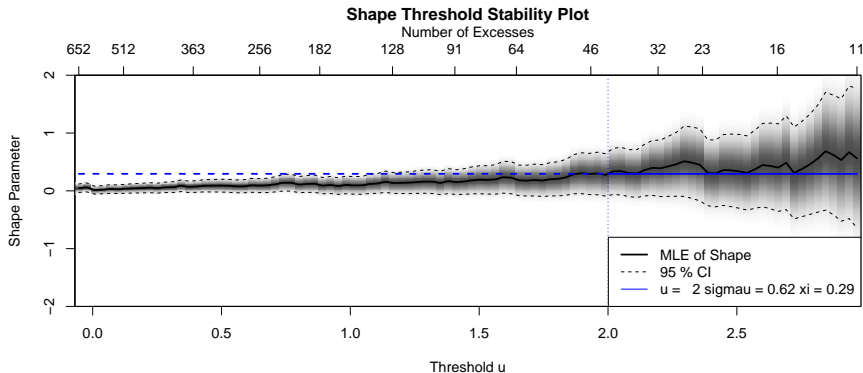
Choosing the threshold: MRL plot

```
> library(evmix)
> mrlplot(dowjones$returns,
+         try.thresh=c(quantile(dowjones$returns,0.9),2),
+         ylim = c(-2,4))
```



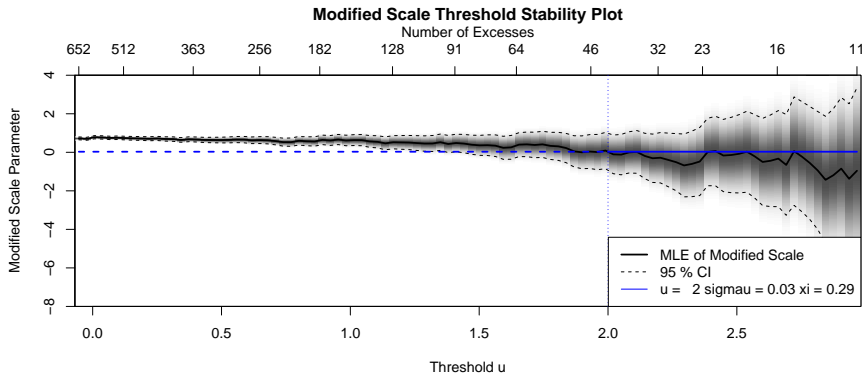
Choosing the threshold: stability plots

```
> tshapeplot(dowjones$returns,ylim=c(-2,2),try.thresh = 2)
```



Choosing the threshold: stability plots

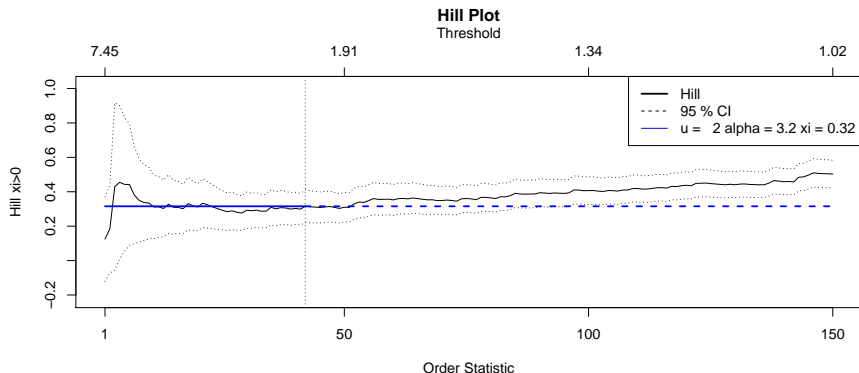
```
> tscaleplot(dowjones$returns,ylim=c(-8,4),try.thresh = 2)
```



Choosing the threshold: hill plot

This is only valid for heavy-tailed distributions $\xi > 0$

```
> hillplot(dowjones$returns, orderl=c(1,150),  
+          try.thresh= 2,xlab="Order Statistic")
```



Fitting the GPD to exceedances

```
> dow_fit <- fevd(dowjones$returns, type="GP", threshold = 2)
> summary(dow_fit)
```

Estimated parameters:

scale	shape
0.6183804	0.2941935

Standard Error Estimates:

scale	shape
0.1496571	0.1918915

Fitting the GPD to exceedances

```
> dow_fit2 <- fevd(dowjones$returns, type="GP",  
+                  threshold=quantile(dowjones$returns, 0.9))  
> summary(dow_fit2)
```

Estimated parameters:


scale	shape
0.6593380	0.1490148

Standard Error Estimates:


scale	shape
0.08408249	0.09380125

Estimated densities

Thresh=2



Tresh=Q(0.9)



Estimated parameters and return levels

Thresh = 2

	lower	mean	upper
scale	0.325	0.618	0.911
shape	-0.081	0.294	0.670

Thresh = $q(0.9)$


	lower	mean	upper
scale	0.494	0.659	0.824
shape	-0.034	0.194	0.333

	lower	mean	upper
2-year	3.43	5.22	7.00
20-year	1.35	10.4	19.40
100-year	-5.74	16.7	39.20


	lower	mean	upper
2-year	3.77	5.22	7.00
20-year	4.25	8.56	12.86
100-year	3.56	11.7	19.97

Quantile plots



Thresh=2



Tresh=Q(0.9)




Are returns independent?



Extremal index


The extremal index θ measures the level of clustering of extreme events ($\theta = 1$ corresponding to independence)



C.A.T. Ferro, J. Segers (2003) Inference for clusters of extreme values . *Journal of the Royal Statistical Society Series B* 65:545-556.

Declustering

We can use the function `decluster` to extract the cluster maxima.



Results after declustering

Thresh.	r	σ	ξ	Ext.Index
2	0	0.62 (0.15)	0.29 (0.19)	0.74
2	5	0.61 (0.15)	0.31 (0.20)	0.76
2	7	0.79 (0.23)	0.24 (0.22)	1
$q(0.9)$	0	0.66 (0.08)	0.15 (0.09)	0.57
$q(0.9)$	4	0.72 (0.10)	0.12 (0.09)	0.70
$q(0.9)$	7	1.10 (0.22)	0.07 (0.14)	1

Heteroscedasticity

- Extremes may also not be identically distributed.
- The scedasis function gives a representation of tail risk.
- To compute this we need to restrict ourselves to the case $\xi > 0$ and constant.
- The `extremis` package gives an easy implementation for the scedasis.

J.H.J Einmahl, L. de Haan, C. Zhou (2016) Statistics of heteroscedastic extremes.
Journal of the Royal Statistical Society Series B 78:31-51.

Testing for constant ξ

Here we implement the following test

- 1 Compute the Hill estimator $\hat{\xi}$ over the full time series using the k highest observations
- 2 Divide the time series into m blocks of equal length
- 3 Compute the Hill estimator $\hat{\xi}_l$ over each block using the $\lfloor k/m \rfloor + 1$ highest observations
- 4 Compute the test statistics

$$T = \frac{1}{m} \sum_{l=1}^m \left(\frac{\hat{\gamma}_l}{\hat{\gamma}} - 1 \right)^2$$

- 5 Compare kT with the quantile $\chi_{m-1}^2(1 - \alpha)$ for a confidence level α

Results for the Dow Jones series

The values k were chosen to match the threshold used till now.

m	k	p-value
4	42	0.99
4	130	0.67
3	42	0.87
3	130	0.99

The assumption of constant ξ looks tenable.

Fitting the scedasis function

We can use the function `cdensity` to estimate the scedasis.




Figure: Scedasis function: threshold=2 (left), threshold= $q(0.9)$ (right)

Fitting the scedasis cdf





Figure: Scedasis cdf: threshold=2 (left), threshold= $q(0.9)$ (right)

Conclusions

- R provide a set of packages to carry out a large number of inferential routines for extremes, including non-stationary, heteroscedastic extensions.
- Tomorrow we will look at the implementation of multivariate methods
- On Saturday you will actively fit extreme models to data.